

Container essentials

DIY containers

Who am I?

Why this talk?

fun - be curious - know your tools - improve yourself

Containers

- What is a container?
- History
 - since 2002, first patch in 2.4.19 tree
 - main work in the 2.6 series in the kernel

The container is a lie

- The linux kernel has no idea what a container is
- Namespaces - limiting the view for a process
- Cgroups - limiting the resources for a process

So, what is a container

- A process
- In a separate environment

Running a process

```
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    execvp(argv[1], &argv[1]);
    return 0;
}
```

fork, clone and exec

clone - exec - wait

- First we clone, so we have a child thread
- Then, we replace the child thread with the command we wish to run
- Then we wait for the child to return

```
static int child_exec(void *argv)
{
    execvp(argv[0], argv);
}

int main(int argc, char **argv)
{
    int flags = SIGCHLD;

    pid = clone(child_exec, ..., flags, &argv);

    waitpid(pid, NULL, 0)
}
```

Full source - part 1

```
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <errno.h>
#include <string.h>
#include <sched.h>
#include <sys/wait.h>

#define STACKSIZE (1024*1024)
static char child_stack[STACKSIZE];

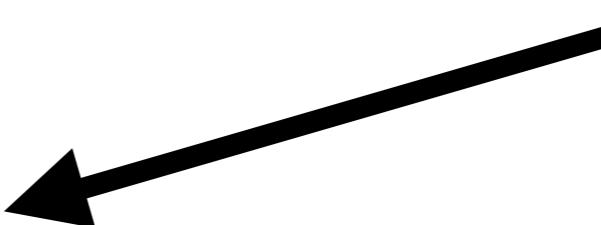
struct clone_args {
    char **argv;
};

// child_exec is the func that will be executed as the result of clone
static int child_exec(void *stuff)
{
    struct clone_args *args = (struct clone_args *)stuff;
    if (execvp(args->argv[0], args->argv) != 0) {
        fprintf(stderr, "failed to execvp arguments %s\n",
                strerror(errno));
        exit(-1);
    }
    // we should never reach here!
    exit(EXIT_FAILURE);
}

// continued in next slide !!!
```

Full source - part 2

```
int main(int argc, char **argv)
{
    struct clone_args args;
    args.argv = &argv[1];

    int clone_flags = SIGCHLD;  Add the flags here

    // the result of this call is that our child_exec will be run in another
    // process returning it's pid
    pid_t pid =
        clone(child_exec, child_stack + STACKSIZE, clone_flags, &args);
    if (pid < 0) {
        fprintf(stderr, "clone failed !!!! %s\n", strerror(errno));
        exit(EXIT_FAILURE);
    }
    // lets wait on our child process here before we, the parent, exits
    if (waitpid(pid, NULL, 0) == -1) {
        fprintf(stderr, "failed to wait pid %d\n", pid);
        exit(EXIT_FAILURE);
    }
    exit(EXIT_SUCCESS);
}
```

Running our container

```
gcc -o container container.c  
./container echo "I'm inside a container!"
```

Namespaces

Enables a process to have different views of the system
than other processes.

it alters the view of the system for a process

Existing namespaces

- CLONE_NEWUTS: Hostname and NIS domain name
- CLONE_NEWPID: Process IDs
- CLONE_NEWNET: Network devices, stacks, ports,...
- CLONE_NEWNS: Mount points
- CLONE_NEWUSER: User and group IDs
- CLONE_NEWIPC: SysV IPC, POSIX message queues

CLONE_NEWUTS

./containers /bin/hostname bar

```
int clone_flags = SIGCHLD | CLONE_NEWUTS;
```

These are flags, so we logical-OR them

CLONE_NEWNET

./containers /bin/ip a

CLONE_NEWNS

mount / umount

provide a new root through bind-(re)mounting

./container /bin/bash

umount /proc - recheck in the main box

CLONE_NEWPID

bin/containers /bin/ps
we need to re-mount proc

Remounting /proc

```
#include <sys/mount.h>

// child_exec is the func that will be executed as the result of clone
static int child_exec(void *stuff)
{
    struct clone_args *args = (struct clone_args *)stuff;

    if (umount("/proc") != 0) {
        fprintf(stderr, "failed umount /proc %s\n",
                strerror(errno));
        exit(-1);
    }
    if (mount("proc", "/proc", "proc", 0, "") != 0) {
        fprintf(stderr, "failed mount /proc %s\n",
                strerror(errno));
        exit(-1);
    }
}
```

unshare

We just created unshare

```
ubuntu@ubuntu-xenial:~$ unshare -h
```

Usage:

```
unshare [options] <program> [<argument>...]
```

Run a program with some namespaces unshared from the parent.

Options:

-m, --mount [=<file>]	unshare mounts namespace
-u, --uts [=<file>]	unshare UTS namespace (hostname etc)
-i, --ipc [=<file>]	unshare System V IPC namespace
-n, --net [=<file>]	unshare network namespace
-p, --pid [=<file>]	unshare pid namespace
-U, --user [=<file>]	unshare user namespace
-f, --fork	fork before launching <program>
--mount-proc [=<dir>]	mount proc filesystem first (implies --mount)
-r, --map-root-user	map current user to root (implies --user)
--propagation slave shared private unchanged	modify mount propagation in mount namespace
-s, --setgroups allow deny	control the setgroups syscall in user ns
-h, --help	display this help and exit
-V, --version	output version information and exit

For more details see `unshare(1)`.

Cgroups

Controls the resources available to a process/processes
Used to control memory / cpu / ... resources